

深度學習專題:水果分類

成員:林家平

引用資料集

功能表

+

創造

探 家

比賽

數據

法 法典

評 討論

學 學習

更多

您的工作

最近查看

CNN的水果和蔬菜分類

查看活動事件

搜索

米哈伊·奧雷爾安 · 一年前更新

2582

新筆記本

下載 (1 GB)

水果360

包含 131 種水果和蔬菜的 90380 張圖像的數據集

數據卡 代碼 (471) 討論 (24)

關於數據集

水果 360 數據集：包含水果和蔬菜的圖像數據集

版本：2020.05.18.0

可用性 資訊
8.75

許可證
CC BY-SA 4.0

預期更新頻率



算法:CNN

```
class FruitCnnModel(nn.Module):
    def __init__(self):
        super().__init__()
        self.network = nn.Sequential(

            nn.Conv2d(3, 16, kernel_size=2, padding=1), #99
            nn.BatchNorm2d(16),
            nn.ReLU(),
            nn.MaxPool2d(2, 2), #99/2=50

            nn.Conv2d(16, 32, kernel_size=2, stride=1, padding=1), #49
            nn.BatchNorm2d(32),
            nn.ReLU(),
            nn.MaxPool2d(2, 2), #49/2=25

            nn.Conv2d(32, 64, kernel_size=2, stride=1, padding=1), #24
            nn.BatchNorm2d(64),
            nn.ReLU(),
            nn.MaxPool2d(5, 5), #24/5=5

            nn.Flatten(),
            nn.Dropout(0.5),
            nn.ReLU(),
            nn.Linear(64*5*5, 131))

    def forward(self, xb):
        return self.network(xb)
```

訓練

```
criterion = nn.CrossEntropyLoss()
optimizer = optim.Adam(model.parameters(), lr=0.001)
n_epochs = 4
valid_loss_min = np.Inf
```

```
running epoch: 1
    Training Loss: 1.306246      Validation Loss: 0.146474
Validation loss decreased (inf --> 0.146474). Saving model ...
running epoch: 2
    Training Loss: 0.137370      Validation Loss: 0.041651
Validation loss decreased (0.146474 --> 0.041651). Saving model ...
running epoch: 3
    Training Loss: 0.060066      Validation Loss: 0.013181
Validation loss decreased (0.041651 --> 0.013181). Saving model ...
running epoch: 4
    Training Loss: 0.035998      Validation Loss: 0.007642
Validation loss decreased (0.013181 --> 0.007642). Saving model ...
```

測試結果

```
def test(loaders, model, criterion, use_cuda):

    test_loss = 0.
    correct = 0.
    total = 0.
    model.eval()
    for batch_idx, (data, target) in enumerate(loaders):

        if use_cuda:
            data, target = data.cuda(), target.cuda()

        output = model(data)

        loss = criterion(output, target)

        test_loss = test_loss + ((1 / (batch_idx + 1)) * (loss.data - test_loss))

        pred = output.data.max(1, keepdim=True)[1]

        correct += np.sum(pred.eq(target.data.view_as(pred)).cpu().numpy())
        total += data.size(0)

    print('Test Loss: {:.6f}'.format(test_loss))
    print('Test Accuracy: %2d%% (%2d/%2d)' % (100. * correct / total, correct, total))


use_cuda = torch.cuda.is_available()
test(test_loader, model, criterion, use_cuda)
```

Test Loss: 0.077323
Test Accuracy: 98% (22305/22688)

損失值曲線圖

```
train_losses = [x for x in train_losses]
epochs=[1,2,3,4]
val_losses = [x for x in valid_losses]
plt.plot(epochs,train_losses, '-bx')
plt.plot(epochs,val_losses, '-rx')
plt.xlabel('epoch')
plt.ylabel('loss')
plt.ylim(0,2)
plt.legend(['Training', 'Validation'])
plt.title('Loss vs. No. of epochs');

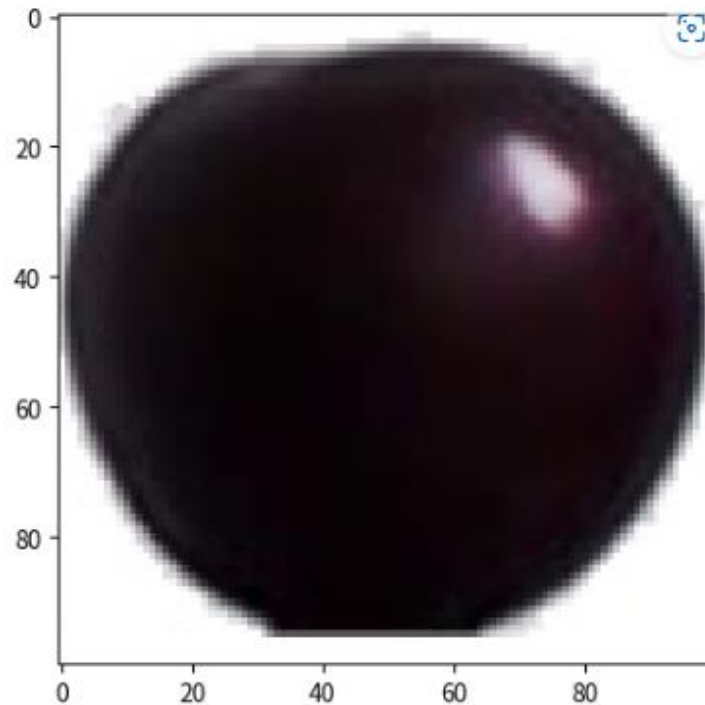
plt.show()
```



模型測試

```
In [11]: img, label = test_data[5000]
plt.imshow(img.permute(1, 2, 0))
print('Label:', train_data.classes[label], 'Predicted:', predict_image(img, model))
```

Label: Cherry Wax Black Predicted: Cherry Wax Black



+ • 蝴蝶和飛蛾圖像分類

 GERRY · UPDATED 3 MONTHS AGO

▲ 86

New Notebook

Download (476 MB)

Butterfly & Moths Image Classification 100 species

12639 train, 500 test, 500 validation images 224 X 224 X 3 jpg format

Data Card

Code (29)

Discussion (1)



+

•

○

+ • ○ 算法:CNN

```
class Butterfly_vs_Classifier_CnnModel(nn.Module):
    def __init__(self):
        super().__init__()
        self.network = nn.Sequential(

            nn.Conv2d(3, 16, kernel_size=5, padding=0),
            nn.BatchNorm2d(16),
            nn.ReLU(),
            nn.MaxPool2d(2, 2),

            nn.Conv2d(16, 32, kernel_size=5, stride=1, padding=0),
            nn.BatchNorm2d(32),
            nn.ReLU(),
            nn.MaxPool2d(2, 2),

            nn.Conv2d(32, 16, kernel_size=3, stride=1, padding=0),
            nn.BatchNorm2d(16),
            nn.ReLU(),
            nn.MaxPool2d(2, 2),

            nn.Conv2d(16, 8, kernel_size=3, stride=1, padding=0),
            nn.BatchNorm2d(8),
            nn.ReLU(),
            nn.MaxPool2d(2, 2),

            nn.Flatten(),
            nn.Dropout(0.5),
            nn.ReLU(),
            nn.Linear(8*11*11, 100))

    def forward(self, xb):
        return self.network(xb)

model = FruitCnnModel()
model
```

+

•

○

+
•
○

訓練

```
criterion = nn.CrossEntropyLoss()  
optimizer = optim.Adam(model.parameters(), lr=0.0001)
```

```
n_epochs = 50
```

```
valid_loss_min = np.Inf
```

```
running epoch: 1  
  Training Loss: 1.219365      Validation Loss: 1.002321  
Validation loss decreased (inf --> 1.002321). Saving model ...  
running epoch: 2  
  Training Loss: 1.191668      Validation Loss: 1.005435  
running epoch: 3  
  Training Loss: 1.182129      Validation Loss: 1.006855  
running epoch: 4  
  Training Loss: 1.167862      Validation Loss: 0.974304  
Validation loss decreased (1.002321 --> 0.974304). Saving model ...  
running epoch: 5  
  Training Loss: 1.146752      Validation Loss: 1.051923  
running epoch: 6  
  Training Loss: 1.145451      Validation Loss: 1.018808  
running epoch: 7  
  Training Loss: 1.119549      Validation Loss: 0.967221  
Validation loss decreased (0.974304 --> 0.967221). Saving model ...  
running epoch: 8  
  Training Loss: 1.100453      Validation Loss: 1.018015  
running epoch: 9  
  Training Loss: 1.095090      Validation Loss: 1.034428  
running epoch: 10  
  Training Loss: 1.080947      Validation Loss: 0.958184  
Validation loss decreased (0.967221 --> 0.958184). Saving model ...  
running epoch: 11  
  Training Loss: 1.087827      Validation Loss: 0.925947  
Validation loss decreased (0.958184 --> 0.925947). Saving model ...  
running epoch: 12  
  Training Loss: 1.055455      Validation Loss: 0.938659  
running epoch: 13  
  Training Loss: 1.058292      Validation Loss: 0.950806  
running epoch: 14  
  Training Loss: 1.041591      Validation Loss: 0.902710  
Validation loss decreased (0.925947 --> 0.902710). Saving model ...  
running epoch: 15  
  Training Loss: 1.031274      Validation Loss: 0.952339  
running epoch: 16  
  Training Loss: 1.024050      Validation Loss: 0.918353  
running epoch: 17  
  Training Loss: 0.984280      Validation Loss: 0.911466  
running epoch: 18  
  Training Loss: 0.997803      Validation Loss: 0.887306  
Validation loss decreased (0.902710 --> 0.887306). Saving model ...  
running epoch: 19  
  Training Loss: 0.991826      Validation Loss: 0.928942  
running epoch: 20  
  Training Loss: 0.979805      Validation Loss: 0.937860  
running epoch: 21  
  Training Loss: 0.972732      Validation Loss: 0.881339  
Validation loss decreased (0.887306 --> 0.881339). Saving model ...  
running epoch: 22  
  Training Loss: 0.967737      Validation Loss: 0.840881  
Validation loss decreased (0.881339 --> 0.840881). Saving model ...  
running epoch: 23  
  Training Loss: 0.955617      Validation Loss: 0.857133  
running epoch: 24  
  Training Loss: 0.941583      Validation Loss: 0.860032  
running epoch: 25  
  Training Loss: 0.939470      Validation Loss: 0.863752  
running epoch: 26  
  Training Loss: 0.936656      Validation Loss: 0.892765
```

```
running epoch: 26  
  Training Loss: 0.936656      Validation Loss: 0.892765  
running epoch: 27  
  Training Loss: 0.912217      Validation Loss: 0.859003  
running epoch: 28  
  Training Loss: 0.913842      Validation Loss: 0.899580  
running epoch: 29  
  Training Loss: 0.922358      Validation Loss: 0.848848  
running epoch: 30  
  Training Loss: 0.905812      Validation Loss: 0.843210  
running epoch: 31  
  Training Loss: 0.889700      Validation Loss: 0.845844  
running epoch: 32  
  Training Loss: 0.881824      Validation Loss: 0.845093  
running epoch: 33  
  Training Loss: 0.879473      Validation Loss: 0.829454  
Validation loss decreased (0.840881 --> 0.829454). Saving model ...  
running epoch: 34  
  Training Loss: 0.894274      Validation Loss: 0.895647  
running epoch: 35  
  Training Loss: 0.877130      Validation Loss: 0.834756  
running epoch: 36  
  Training Loss: 0.875628      Validation Loss: 0.817081  
Validation loss decreased (0.829454 --> 0.817081). Saving model ...  
running epoch: 37  
  Training Loss: 0.883569      Validation Loss: 0.836547  
running epoch: 38  
  Training Loss: 0.863131      Validation Loss: 0.810686  
Validation loss decreased (0.817081 --> 0.810686). Saving model ...  
running epoch: 39  
  Training Loss: 0.848675      Validation Loss: 0.826044  
running epoch: 40  
  Training Loss: 0.851950      Validation Loss: 0.809089  
Validation loss decreased (0.810686 --> 0.809089). Saving model ...  
running epoch: 41  
  Training Loss: 0.844693      Validation Loss: 0.843245  
running epoch: 42  
  Training Loss: 0.852901      Validation Loss: 0.859733  
running epoch: 43  
  Training Loss: 0.829086      Validation Loss: 0.808046  
Validation loss decreased (0.809089 --> 0.808046). Saving model ...  
running epoch: 44  
  Training Loss: 0.850447      Validation Loss: 0.815440  
running epoch: 45  
  Training Loss: 0.824561      Validation Loss: 0.831608  
running epoch: 46  
  Training Loss: 0.820571      Validation Loss: 0.837813  
running epoch: 47  
  Training Loss: 0.814792      Validation Loss: 0.805094  
Validation loss decreased (0.808046 --> 0.805094). Saving model ...  
running epoch: 48  
  Training Loss: 0.828432      Validation Loss: 0.848318  
running epoch: 49  
  Training Loss: 0.801110      Validation Loss: 0.806240  
running epoch: 50  
  Training Loss: 0.804431      Validation Loss: 0.819413  
Test Loss: 0.752096  
Test Accuracy: 80% (403/500)
```

• 測試結果

Test Loss: 0.752096

Test Accuracy: 80% (403/500)

+

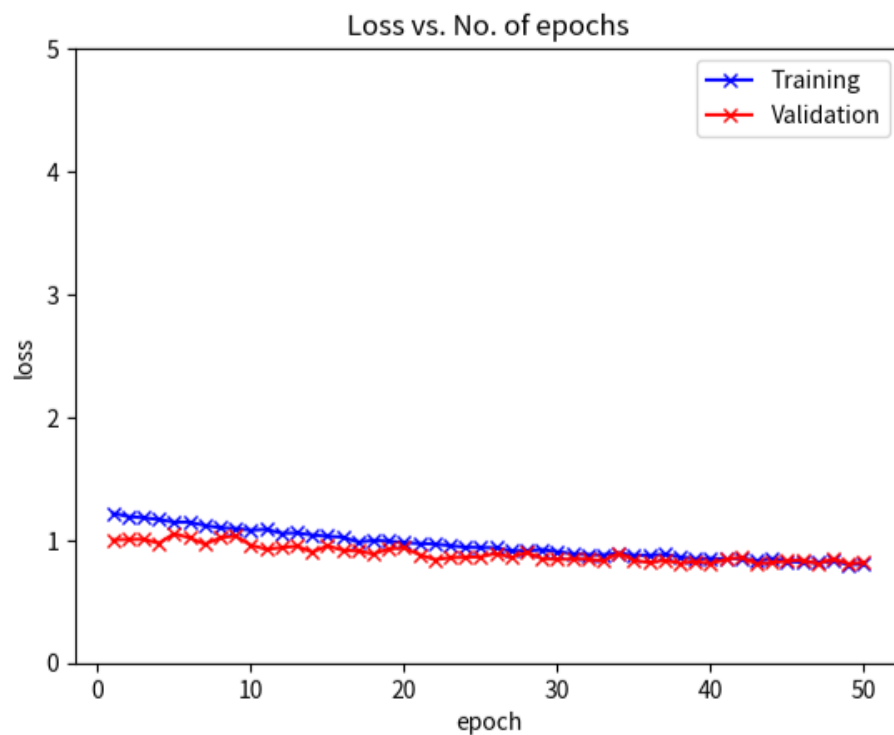
○

•

+ ○ • 損失值曲線圖

```
train_losses = [x for x in train_losses]
epochs=[i for i in range(1,51)]
val_losses = [x for x in valid_losses]
plt.plot(epochs,train_losses, '-bx')
plt.plot(epochs,val_losses, '-rx')
plt.xlabel('epoch')
plt.ylabel('loss')
plt.ylim(0,5)
plt.legend(['Training', 'Validation'])
plt.title('Loss vs. No. of epochs');


plt.show()
```



+ • 模型測試

```
img, label = test_data[20]  
plt.imshow(img.permute(1, 2, 0))  
print('Label:', train_data.classes[label], 'Predicted:', predict_image(img, model))
```

Label: APPOLLO Predicted: APPOLLO





謝謝觀看

